



# Introduction to File Systems

David E. Culler  
CS162 – Operating Systems and Systems  
Programming  
Lecture 5  
Sept 12, 2019

Reading: A&D 5.8, 11.1-2  
HW 1 due 9/18  
Proj 1 Design Doc 9/17



# Objective of this lecture

---

- Resolve tension in understanding Threads
- Show how Operating System functionality distributes across layers in the system.
- Introduce I/O & storage services – i.e., file systems



# Review: Threads

---

- Independently schedulable entity
- Sequential thread of execution that runs concurrently with other threads
  - It can block waiting for something while others progress
  - It can work in parallel with others (ala cs61c)
- Has local state (its stack) and shared (static data and heap)
- In the absence of synchronization operations, arbitrary interleaving of threads may occur




# Recall: Thread State

---

- State shared by all threads in process/addr space
  - Content of memory (global variables, heap)
  - I/O state (file system, network connections, etc)
- Execution Stack (logically private to thread)
  - Parameters, temporary variables
  - Return PCs while called procedures are executing
- State for each thread
  - CPU registers (including, program counter)
  - Ptr to Execution stack
  - Kept in Thread Control Block, when thread not running
- Scheduler works on TCBs




# Recall: Thread Abstraction



- Illusion: Infinite number of processors



# Recall: Thread Abstraction



- Illusion: Infinite number of processors
- Reality: Threads execute with variable “speed”
  - Programs must be designed to work with any schedule



# Recall: Synchronization

---

- **Mutual Exclusion:** Ensuring only one thread does a particular thing at a time (one thread *excludes* the others)
- **Critical Section:** Code exactly one thread can execute at once
  - Result of mutual exclusion
- **Lock:** An object only one thread can hold at a time
  - **Provides** mutual exclusion
- Offers two **atomic** operations:
  - `Lock.Acquire()` – wait until lock is free; then grab
  - `Lock.Release()` – Unlock, wake up waiters
- Need other tools for “cooperation”
  - e.g., Java monitors, semaphores, condition variables)

# Little Example: Stack of Strings (SoS)

---



```
struct str_lst_elem {  
    char *str;  
    struct str_lst_elem *next;  
};  
  
struct str_lst {  
    struct str_lst_elem *head;  
};  
  
void str_lst_init(struct str_lst *lst) {  
    lst->head = NULL;  
};
```



# SoS (cont)



```
void str_lst_push(struct str_lst *lst, char *str) {
    struct str_lst_elem *new_elem = malloc(sizeof(struct str_lst_elem));
    new_elem->str = str;

    new_elem->next = lst->head;
    lst->head = new_elem;

};

char *str_lst_pop(struct str_lst *lst) {
    char *topval;

    struct str_lst_elem *top = lst->head;
    if (!top) {
        topval = NULL;
    } else {
        topval = top->str;
        lst->head = top->next;
    }

    return topval;
};
```

Must be atomic if  
multiple threads

Must be atomic if  
multiple threads



# Thread Safe: Stack of Strings

---

```
struct str_lst_elem {  
    char *str;  
    struct str_lst_elem *next;  
};  
  
struct str_lst {  
    struct str_lst_elem *head;  
    pthread_mutex_t lock;  
};  
  
void str_lst_init(struct str_lst *lst) {  
    lst->head = NULL;  
    pthread_mutex_init(&lst->lock, NULL);  
};
```



# Thread safe: SoS (cont)

```
void str_lst_push(struct str_lst *lst, char *str) {
    struct str_lst_elem *new_elem = malloc(sizeof(struct str_lst_elem));
    new_elem->str = str;
    pthread_mutex_lock (&lst->lock);
    new_elem->next = lst->head;
    lst->head = new_elem;
    pthread_mutex_unlock (&lst->lock);
};

char *str_lst_pop(struct str_lst *lst) {
    char *topval;
    pthread_mutex_lock (&lst->lock);
    struct str_lst_elem *top = lst->head;
    if (!top) {
        topval = NULL;
    } else {
        topval = top->str;
        lst->head = top->next;
    }
    pthread_mutex_unlock (&lst->lock);
    return topval;
};
```

Critical Section

Critical Section

# Lock Implementation Conundrum

---



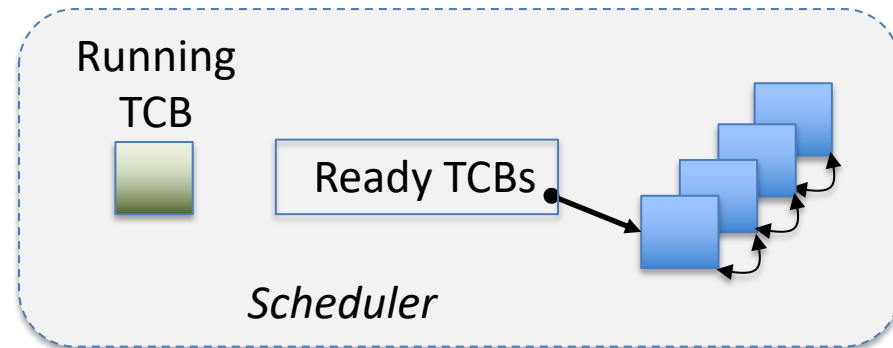

# Lock Implementation Conundrum



- Manipulating a data structure, like a list, requires a sequence of operations that must be atomic
- To make the list thread safe, protect it with a lock
  - Operations on the list are a critical section
  - i.e., lock; manipulate; unlock (*lots of places in the kernel*)
- Lock implementation needs to manipulate lists (of TCBs)
  - Thread that tries to acquire a busy lock is placed on the list of threads waiting on the lock (!!!) – and some other thread scheduled
  - Releasing a lock causes a thread to be removed from the lock's list and placed on the scheduler's list of ready threads
- How do we create critical sections for the lock acquire/release operations themselves ???
- We disable interrupts so no other thread can interleave with this kernel code



# Basic Lock Implementation



```
Acquire(*lock) {  
    disable interrupts;  
    if (lock->value == BUSY) {  
        put thread on lock's wait_Q  
        "i.e, Go to sleep"  
        allow a ready thread to run  
    } else {  
        lock->value = BUSY;  
    }  
    enable interrupts;  
}
```

```
Release(*lock) {  
    disable interrupts;  
    if (any TCB on lock wait_Q) {  
        "i.e., lock busy";  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        lock->value = FREE;  
    }  
    enable interrupts;  
}
```

# Is that all?

---




- There are further subtleties about when interrupts are re-enabled
  - We'll tackle this later, as we get closer to Project 2
- The low level mechanics of thread switch are so simple and subtle that it may still seem like magic
- Still a question of how much of “threads” and “synchronization” could be moved out of the kernel to user level (for performance).
  - later



# Recall: Semaphores

- No negative values
- Only operations allowed are P and V
  - can't read or write value, except to set it initially
- Operations must be atomic
  - Two P's together can't decrement value below zero
  - Similarly, thread going to sleep in P won't miss wakeup from V – even if they both happen at same time
- Semaphore from railway analogy
  - Here is a semaphore initialized to 2 for resource control:





# Recall: Important Semaphore Patterns



- **Mutual Exclusion:** (Like lock)

- Called a "binary semaphore"

- `initial value of semaphore = 1;`

- `semaphore.down();`

- `// Critical section goes here`

- `semaphore.up();`

- **Signaling other threads, e.g. ThreadJoin**

- `Initial value of semaphore = 0`

- ```
ThreadJoin {  
    semaphore.down();  
}
```

- ```
ThreadFinish {  
    semaphore.up();  
}
```



# Intuition for Semaphores

---

- What do you need to wait for?
  - Example: Critical section to be finished
  - Example: Queue to be non-empty, or no longer full
  - Example: Some thread to be done with something
- What can you count that will be 0 when you need to wait?
  - Example: # of threads currently in critical section
  - Example: # of items currently in queue
  - Example: # of free slots in array
  - Example: status of 1 for still active
- Can use semaphore operations to maintain count



# So what's in our PCB now?

---

- Process ID, name, etc
- Thread object(s) – TCBs
  - Place to save registers when not running
  - Thread status
  - Links to form lists
- Thread Stack
- Lock object for *any lock used by its kernel thread*
- User level lock info ???

# Processes and Threads

---



- ???








# Processes issue syscalls ...

---

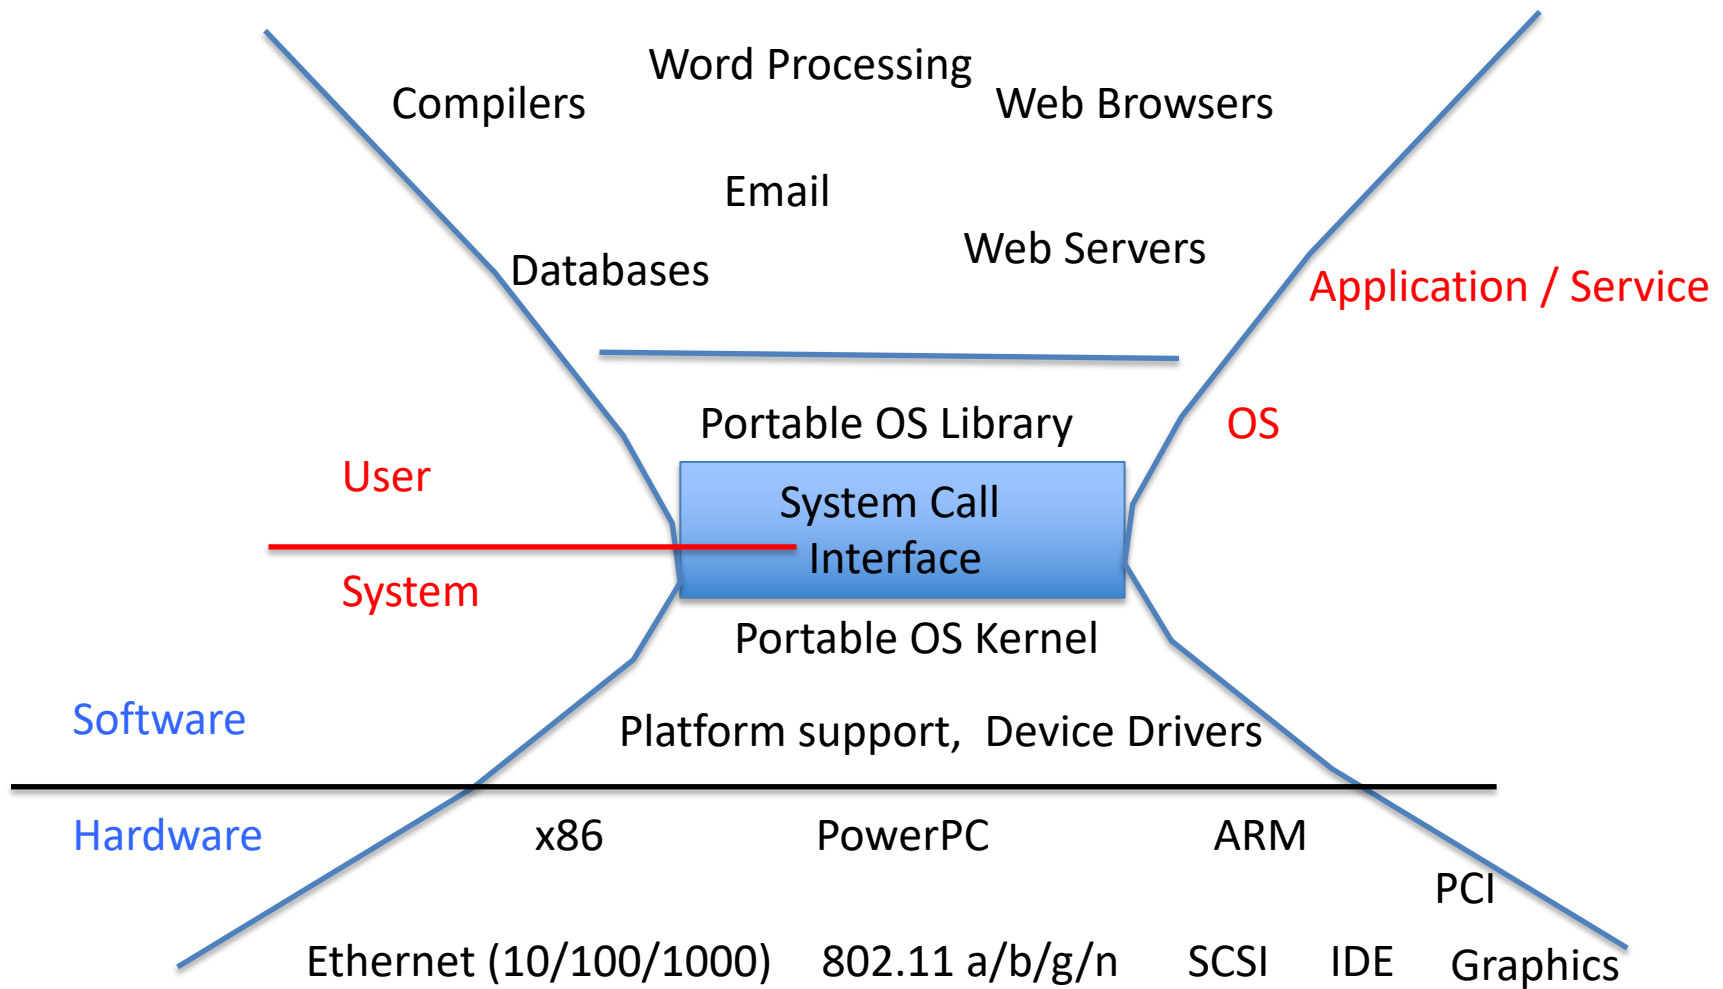
- You said that “applications request services from the operating system via ***syscall***, but ...”
- I’ve been writing all sort of useful applications and I never ever saw a “syscall” !!!
- That’s right.
- It was buried in the programming language runtime library (e.g., libc.a)
- ... Layering

# OS run-time library





# Recall: A Kind of Narrow Waist





# POSIX I/O: Everything is a “File”

---

Identical interface for:

- Devices (terminals, printers, etc.)
- Regular files on disk
- Networking (sockets)
- Local interprocess communication (pipes, sockets)

Based on **open()**, **read()**, **write()**, and **close()**





# POSIX I/O Design Patterns

---

- Open before use
  - Access control check, setup happens here
- Byte-oriented
  - Least common denominator
  - OS responsible for hiding the fact that real devices may not work this way (e.g. hard drive stores data in blocks)
- Explicit close




# POSIX I/O: Kernel Buffering

---

- Reads are buffered
  - Part of making everything byte-oriented
  - Process is **blocked** while waiting for device
  - Let other processes run while gathering result
- Writes are buffered
  - Complete in background (more later on)
  - Return to user when data is “handed off” to kernel

# I/O & Storage Layers

Application / Service



*streams*


*handles*

*registers*

*descriptors*

*Commands and Data Transfers*

*Disks, Flash, Controllers, DMA*





# The file system abstraction


---

- File
  - Named collection of data in a file system
  - POSIX File data: sequence of bytes
    - Could be text, binary, serialized objects, ...
  - File Metadata: information about the file
    - Size, Modification Time, Owner, Security info
    - Basis for access control
- Directory
  - “Folder” containing files & Directories
  - Hierarchical (graphical) naming
    - Path through the directory graph
    - Uniquely identifies a file or directory
      - /home/ff/cs162/public\_html/fa14/index.html
  - Links and Volumes (later)

# C high level File API – streams (review)



- Operate on “streams” - sequence of bytes, whether text or data, with a position



```
#include <stdio.h>
FILE *fopen( const char *filename, const char *mode );
int fclose( FILE *fp );
```

Mode Text	Binary	Descriptions
r	rb	Open existing file for reading
w	wb	Open for writing; created if does not exist
a	ab	Open for appending; created if does not exist
r+	rb+	Open existing file for reading & writing.
w+	wb+	Open for reading & writing; truncated to zero if exists, create otherwise
a+	ab+	Open for reading & writing. Created if does not exist. Read from beginning, write as append

Don't forget to flush

# Connecting Processes, Filesystem, and Users

---



- Process has a ‘current working directory’
- Absolute Paths
  - /home/oski/cs162
- Relative paths
  - index.html, ./index.html - current WD
  - ../index.html - parent of current WD
  - ~, ~cs162 - home directory

# C API Standard Streams — `stdio.h`

---



- Three predefined streams are opened implicitly when the program is executed.
  - `FILE *stdin` – normal source of input, can be redirected
  - `FILE *stdout` – normal source of output, can too
  - `FILE *stderr` – diagnostics and errors
- `STDIN` / `STDOUT` enable composition in Unix
- All can be redirected
  - `cat hello.txt | grep "World!"`
  - **cat's `stdout`** goes to **grep's `stdin`**



# C high level File API – stream ops

```
#include <stdio.h>

// character oriented
int fputc( int c, FILE *fp );           // rtn c or EOF on err
int fputs( const char *s, FILE *fp );   // rtn >0 or EOF

int fgetc( FILE * fp );
char *fgets( char *buf, int n, FILE *fp );

// block oriented
size_t fread(void *ptr, size_t size_of_elements,
              size_t number_of_elements, FILE *a_file);

size_t fwrite(const void *ptr, size_t size_of_elements,
              size_t number_of_elements, FILE *a_file);

// formatted
int fprintf(FILE *restrict stream, const char *restrict
format, ...);
int fscanf(FILE *restrict stream, const char *restrict format,
... );
```





# C Streams: char by char I/O

---

```
#include <stdio.h>
```

```
int main(void) {  
    FILE* input = fopen("input.txt", "r");  
    FILE* output = fopen("output.txt", "w");  
    int c;  
  
    c = fgetc(input);  
    while (c != EOF) {  
        fputc(output, c);  
        c = fgetc(input);  
    }  
    fclose(input);  
    fclose(output);  
}
```



# What if we wanted block by block I/O?

```
#include <stdio.h>
// character oriented
int fputc(int c, FILE *fp);           // rtn c or EOF on err
int fputs(const char *s, FILE *fp);   // rtn >0 or EOF

int fgetc( FILE * fp );
char *fgets( char *buf, int n, FILE *fp );
```

```
// block oriented
size_t fread(void *ptr, size_t size_of_elements,
              size_t number_of_elements, FILE *a_file);

size_t fwrite(const void *ptr, size_t size_of_elements,
              size_t number_of_elements, FILE *a_file);
```

```
// formatted
int fprintf(FILE *restrict stream, const char *restrict format, ...);
int fscanf(FILE *restrict stream, const char *restrict format, ...);
```



# stdio Block-by-Block I/O

---

```
#include <stdio.h>
#define BUFFER_SIZE 1024
int main(void) {
    FILE* input = fopen("input.txt", "r");
    FILE* output = fopen("output.txt", "w");
    char buffer[BUFFER_SIZE];
    size_t length;
    length = fread(buffer, BUFFER_SIZE, sizeof(char), input);

    }
}
```



# stdio Block-by-Block I/O

---

```
#include <stdio.h>
#define BUFFER_SIZE 1024
int main(void) {
    FILE* input = fopen("input.txt", "r");
    FILE* output = fopen("output.txt", "w");
    char buffer[BUFFER_SIZE];
    size_t length;
    length = fread(buffer, BUFFER_SIZE, sizeof(char), input);
    while (length > 0) {
        fwrite(buffer, length, sizeof(char), output);
        length = fread(buffer, BUFFER_SIZE, sizeof(char), input);
    }
    fclose(input);
    fclose(output);
}
```



# Aside: Systems Programming

---

- Systems programmers are paranoid
- We should really be writing things like:

```
FILE* input = fopen("input.txt", "r");  
if (input == NULL) {  
    // Prints our string and error msg.  
    perror("Failed to open input file")  
}
```
- Be **thorough** about checking return values
  - Want failures to be systematically caught and dealt with





# C Stream API: Positioning

```
int fseek(FILE *stream, long int offset, int whence);
```

```
long int ftell (FILE *stream)
```

```
void rewind (FILE *stream)
```



- Preserves high level abstraction of a uniform stream of objects

# What's below the surface ??



Application / Service

High Level I/O

*streams*

Low Level I/O

*handles*

Syscall

*registers*

File System


*descriptors*

I/O Driver

*Commands and Data Transfers*



*Disks, Flash, Controllers, DMA*





# C Low level I/O

- Operations on File Descriptors – as OS object representing the state of a file
  - User has a “handle” on the descriptor

```
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>

int open (const char *filename, int flags [, mode_t mode])
int creat (const char *filename, mode_t mode)
int close (int filedes)
```

Bit vector of:

- Access modes (Rd, Wr, ...)
- Open Flags (Create, ...)
- Operating modes (Appends, ...)

Bit vector of Permission Bits:

- User | Group | Other X R | W | X

[http://www.gnu.org/software/libc/manual/html\\_node/Opening-and-Closing-Files.html](http://www.gnu.org/software/libc/manual/html_node/Opening-and-Closing-Files.html)



# C Low Level: standard descriptors

---



```
#include <unistd.h>
```

```
STDIN_FILENO - macro has value 0
```

```
STDOUT_FILENO - macro has value 1
```

```
STDERR_FILENO - macro has value 2
```

```
int fileno (FILE *stream)
```

```
FILE * fdopen (int filedes, const char *opentype)
```

- Crossing levels: File descriptors vs. streams
- Don't mix them!



# C Low Level Operations

---

`ssize_t read (int filedes, void *buffer, size_t maxsize)`

- returns bytes read, 0 => EOF, -1 => error

`ssize_t write (int filedes, const void *buffer, size_t size)`

- returns bytes written

`off_t lseek (int filedes, off_t offset, int whence)`

`int fsync (int fildes) – wait for i/o to finish`

`void sync (void) – wait for ALL to finish`

- When write returns, data is on its way to disk and can be read, but it may not actually be permanent!



# A little example: lowio.c

---

```
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>

int main() {
    char buf[1000];
    int      fd = open("lowio.c", O_RDONLY, S_IRUSR | S_IWUSR);
    ssize_t  rd = read(fd, buf, sizeof(buf));
    int      err = close(fd);
    ssize_t  wr = write(STDOUT_FILENO, buf, rd);
}
```



# And lots more !

---

- TTYs versus files
- Memory mapped files
- File Locking
- Asynchronous I/O
- Generic I/O Control Operations
- Duplicating descriptors

```
int dup2 (int old, int new)
int dup (int old)
```

# Another: lowio-std.c



```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>

#define BUFSIZE 1024

int main(int argc, char *argv[])
{
    char buf[BUFSIZE];
    ssize_t writelen = write(STDOUT_FILENO, "I am a process.\n", 16);

    ssize_t readlen  = read(STDIN_FILENO, buf, BUFSIZE);

    ssize_t strlen    = snprintf(buf, BUFSIZE, "Got %zd chars\n", readlen);

    writelen = strlen < BUFSIZE ? strlen : BUFSIZE;
    write(STDOUT_FILENO, buf, writelen);

    exit(0);
}
```



# Low-Level I/O: Example

---

```
#include <fcntl.h>
#include <unistd.h>

#define BUFFER_SIZE 1024

int main(void) {
    int input_fd = open("input.txt", O_RDONLY);
    int output_fd = open("output.txt", O_WRONLY);
    char buffer[BUFFER_SIZE];
    ssize_t length;
    length = read(input_fd, buffer, BUFFER_SIZE);
    while (length > 0) {
        write(output_fd, buffer, length);
        length = read(input_fd, buffer, BUFFER_SIZE);
    }
    close(input_fd);
    close(output_fd);
}
```



# Low-Level I/O: Other Operations

---

- Operations specific to terminals, devices, networking, ...
- Duplicating descriptors
  - `int dup2(int old, int new);`
  - `int dup(int old);`
- Pipes – bi-directional channel
  - `int pipe(int fileds[2]);`
  - Writes to `fileds[1]` read from `fileds[0]`
- File Locking
- Memory-Mapping Files
- Asynchronous I/O



# Little pipe example

---

```
#include <unistd.h>

#define BUFSIZE 1024
enum PipeSel {rd_pipe = 0, wt_pipe = 1};

int main(int argc, char *argv[])
{
    char *msg = "Message in a pipe.\n";
    char buf[BUFSIZE];
    int pipe_fd[2];
    if (pipe(pipe_fd)) {
        fprintf (stderr, "Pipe failed.\n"); return EXIT_FAILURE;
    }
    ssize_t writelen = write(pipe_fd[wt_pipe], msg, strlen(msg)+1);
    printf("Sent: %s [%ld, %ld]\n", msg, strlen(msg)+1, writelen);

    ssize_t readlen = read(pipe_fd[rd_pipe], buf, BUFSIZE);
    printf("Rcvd: %s [%ld]\n", msg, readlen);
    close(pipe_fd[wt_pipe]);
    close(pipe_fd[rd_pipe]);
}
```



# Inter-Process Communication (IPC)



- One process reads a file the other writes, or ...

```
pid_t pid = fork();
if (pid < 0) {
    fprintf (stderr, "Fork failed.\n");
    return EXIT_FAILURE;
}
if (pid != 0) {
    ssize_t writelen = write(pipe_fd[wt_pipe], msg, msglen);
    printf("Parent: %s [%ld, %ld]\n", msg, msglen, writelen);
    close(pipe_fd[wt_pipe]);
} else {
    ssize_t readlen = read(pipe_fd[rd_pipe], buf, BUFSIZE);
    printf("Child Rcvd: %s [%ld]\n", msg, readlen);
    close(pipe_fd[rd_pipe]);
}
```



# Streams vs. File Descriptors

---

- Streams are **buffered in user memory**:

```
printf("Beginning of line ");  
sleep(10); // sleep for 10 seconds  
printf("and end of line\n");
```

Prints out **everything at once**

- Operations on file descriptors are **visible immediately**

```
write(STDOUT_FILENO, "Beginning of line ", 18);  
sleep(10);  
write("and end of line \n", 16);
```

Outputs "Beginning of line" 10 seconds earlier

# Why Buffer in Userspace? Overhead!

---



- Avoid system call overhead
  - Time to copy registers, transition to kernel mode, jump to system call handler, etc.
- Minimum syscall time: ~100s of nanoseconds
  - Read/write a file byte by byte?
  - Max throughput of **~10MB/second**
  - With **fgetc**? Keeps up with your SSD

# Why Buffer in Userspace? Functionality.

---



- System call operations less capable
  - Simplifies operating system
- Example: No "read until new line" operation
  - Solution: Make a big read syscall, find first new line in userspace
  - Could simulate by one syscall per character, but we already know this is a bad idea



# Key Unix I/O Design Concepts


---

- Uniformity – everything is a file
  - file operations, device I/O, and interprocess communication through open, read/write, close
  - Allows simple composition of programs
    - `find | grep | wc ...`
- Open before use
  - Provides opportunity for access control and arbitration
  - Sets up the underlying machinery, i.e., data structures
- Byte-oriented
  - Even if blocks are transferred, addressing is in bytes
- Kernel buffered reads
  - Streaming and block devices looks the same, read blocks yielding processor to other task
- Kernel buffered writes
  - Completion of out-going transfer decoupled from the application, allowing it to continue
- Explicit close

# What's below the surface ??



Application / Service



*streams*


*handles*

*registers*

*descriptors*

*Commands and Data Transfers*

*Disks, Flash, Controllers, DMA*



# Break

---



# Recall: SYSCALL



syscalls.kernelgrok.com

BCal UCB CS162 cullermayeno Wikipedia Yahoo! News Popular Imported From Safari

## Linux Syscall Reference

Show 10 entries Search:

#	Name	Registers						Definition
		eax	ebx	ecx	edx	esi	edi	
0	<b>sys_restart_syscall</b>	0x00	-	-	-	-	-	kernel/signal.c:2058
1	<b>sys_exit</b>	0x01	int error_code	-	-	-	-	kernel/exit.c:1046
2	<b>sys_fork</b>	0x02	struct pt_regs *	-	-	-	-	arch/alpha/kernel/entry.S:716
3	<b>sys_read</b>	0x03	unsigned int fd	char __user *buf	size_t count	-	-	fs/read_write.c:391
4	<b>sys_write</b>	0x04	unsigned int fd	const char __user *buf	size_t count	-	-	fs/read_write.c:408
5	<b>sys_open</b>	0x05	const char __user *filename	int flags	int mode	-	-	fs/open.c:900
6	<b>sys_close</b>	0x06	unsigned int fd	-	-	-	-	fs/open.c:969
7	<b>sys_waitpid</b>	0x07	pid_t pid	int __user *stat_addr	int options	-	-	kernel/exit.c:1771
8	<b>sys_creat</b>	0x08	const char __user *pathname	int mode	-	-	-	fs/open.c:933
9	<b>sys_link</b>	0x09	const char __user *oldname	const char __user *newname	-	-	-	fs/namei.c:2520

Showing 1 to 10 of 338 entries First Previous 1 2 3 4 5 Next Last

Generated from Linux kernel 2.6.35.4 using **Exuberant Ctags, Python, and DataTables.**  
Project on [GitHub](#). Hosted on [GitHub Pages](#).

- Low level lib parameters are set up in registers and syscall instruction is issued
  - A type of synchronous exception that enters well-defined entry points into kernel





# What's below the surface ??

File descriptor number  
- an int

File Descriptors  
• a struct with all the info  
about the files

Application / Service

High Level I/O

streams

Low Level I/O

handles

Syscall


registers

File System

descriptors

I/O Driver

Commands and Data Transfers  
Disks, Flash, Controllers, DMA





# Internal OS File Descriptor

- Internal Data Structure describing everything about the file
  - Where it resides
  - Its status
  - How to access it

```
lxr.free-electrons.com/source/include/linux/fs.h#L747
746
747 struct file {
748     union {
749         struct llist_node    fu_llist;
750         struct rcu_head      fu_rcuhead;
751     } f_u;
752     struct path              f_path;
753 #define f_dentry              f_path.dentry
754     struct inode              *f_inode; /* caci
755     const struct file_operations *f_op;
756
757     /*
758      * Protects f_ep_links, f_flags.
759      * Must not be taken from IRQ context.
760      */
761     spinlock_t                f_lock;
762     atomic_long_t              f_count;
763     unsigned int               f_flags;
764     fmode_t                    f_mode;
765     struct mutex               f_pos_lock;
766     loff_t                     f_pos;
767     struct fown_struct          f_owner;
768     const struct cred           *f_cred;
769     struct file_ra_state        f_ra;
770
771     u64                         f_version;
772 #ifdef CONFIG_SECURITY
773     void                        *f_security;
774 #endif
775     /* needed for tty driver, and maybe others */
776     void                        *private_data;
777
778 #ifdef CONFIG_EPOLL
779     /* Used by fs/eventpoll.c to link all the hook:
780     struct list_head            f_ep_links;
781     struct list_head            f_tfile_llink;
782 #endif /* #ifdef CONFIG_EPOLL */
783     struct address_space        *f_mapping;
784 } __attribute__((aligned(4))); /* lest something weird
785
```



# File System: from syscall to driver

In fs/read\_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```




# Layer by layer

User App

User library

```
length = read(input_fd, buffer, BUFFER_SIZE);
```

Application / Service



```
ssize_t read(int, void *, size_t){  
    marshal args into registers  
    issue syscall  
    register result of syscall to rtn value  
};
```

Exception  $U \rightarrow K$ , interrupt processing

```
Void syscall_handler (struct intr_frame *f) {  
    unmarshall call#, args from regs  
    dispatch : handlers[call#](args)  
    marshal results fo syscall ret  
}
```

```
ssize_t vfs_read(struct file *file, char  
__user *buf, size_t count, loff_t *pos)  
{  
    UserProcess/File System relationship  
    call device driver to do the work  
}
```

Device Driver



# Low Level Driver

- Associated with particular hardware device
- Registers / Unregisters itself with the kernel
- Handler functions for each of the file operations

```
struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
    ssize_t (*aio_read) (struct kiocb *, const struct iovec *, unsigned long, loff_t);
    ssize_t (*aio_write) (struct kiocb *, const struct iovec *, unsigned long, loff_t);
    int (*readdir) (struct file *, void *, filldir_t);
    unsigned int (*poll) (struct file *, struct poll_table_struct *);
    int (*ioctl) (struct inode *, struct file *, unsigned int, unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *, fl_owner_t id);
    int (*release) (struct inode *, struct file *);
    int (*fsync) (struct file *, struct dentry *, int datasync);
    int (*fasync) (int, struct file *, int);
    int (*flock) (struct file *, int, struct file_lock *);
    [...]
};
```



# File System: from syscall to driver

## In fs/read\_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return 0;
    if (!file->f_op || (!file->f_op->read &&
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count)))
        return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

- Read up to “count” bytes from “file” starting from “pos” into “buf”.
- Return error or number of bytes read.



# File System: from syscall to driver

## In fs/read\_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

Make sure we are  
allowed to read  
this file



# File System: from syscall to driver

## In fs/read\_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

Check if file has  
read methods





# File System: from syscall to driver

## In fs/read\_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count);
        else
            ret = do_sync_read(file, buf, count);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

- Check whether we can write to buf (e.g., buf is in the user space range)
- unlikely(): hint to branch prediction this condition is unlikely



# File System: from syscall to driver

## In fs/read\_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

Check whether we read from a valid range in the file.



# File System: from syscall to driver

## In fs/read\_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

If driver provide a read function (f\_op->read) use it; otherwise use do\_sync\_read()



# File System: from syscall to driver

## In fs/read\_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

Notify the parent of this file that the file was read (see <http://www.fieldses.org/~bfields/kernel/vfs.txt>)



# File System: from syscall to driver

## In fs/read\_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

Update the number of bytes read by “current” task (for scheduling purposes)



# File System: from syscall to driver

## In fs/read\_write.c

```
ssize_t vfs_read(struct file *file, char __user *buf, size_t count, loff_t *pos)
{
    ssize_t ret;
    if (!(file->f_mode & FMODE_READ)) return -EBADF;
    if (!file->f_op || (!file->f_op->read && !file->f_op->aio_read))
        return -EINVAL;
    if (unlikely(!access_ok(VERIFY_WRITE, buf, count))) return -EFAULT;
    ret = rw_verify_area(READ, file, pos, count);
    if (ret >= 0) {
        count = ret;
        if (file->f_op->read)
            ret = file->f_op->read(file, buf, count, pos);
        else
            ret = do_sync_read(file, buf, count, pos);
        if (ret > 0) {
            fsnotify_access(file->f_path.dentry);
            add_rchar(current, ret);
        }
        inc_syscr(current);
    }
    return ret;
}
```

Update the number of read syscalls by “current” task (for scheduling purposes)



# Lower Level Driver

- Associated with particular hardware device
- Registers / Unregisters itself with the kernel
- Handler functions for each of the file operations

```
struct file_operations {
    struct module *owner;
    loff_t (*llseek) (struct file *, loff_t, int);
    ssize_t (*read) (struct file *, char __user *, size_t, loff_t *);
    ssize_t (*write) (struct file *, const char __user *, size_t, loff_t *);
    ssize_t (*aio_read) (struct kiocb *, const struct iovec *, unsigned long, loff_t);
    ssize_t (*aio_write) (struct kiocb *, const struct iovec *, unsigned long, loff_t);
    int (*readdir) (struct file *, void *, filldir_t);
    unsigned int (*poll) (struct file *, struct poll_table_struct *);
    int (*ioctl) (struct inode *, struct file *, unsigned int, unsigned long);
    int (*mmap) (struct file *, struct vm_area_struct *);
    int (*open) (struct inode *, struct file *);
    int (*flush) (struct file *, fl_owner_t id);
    int (*release) (struct inode *, struct file *);
    int (*fsync) (struct file *, struct dentry *, int datasync);
    int (*fasync) (int, struct file *, int);
    int (*flock) (struct file *, int, struct file_lock *);
    [...]
};
```



# Device Drivers

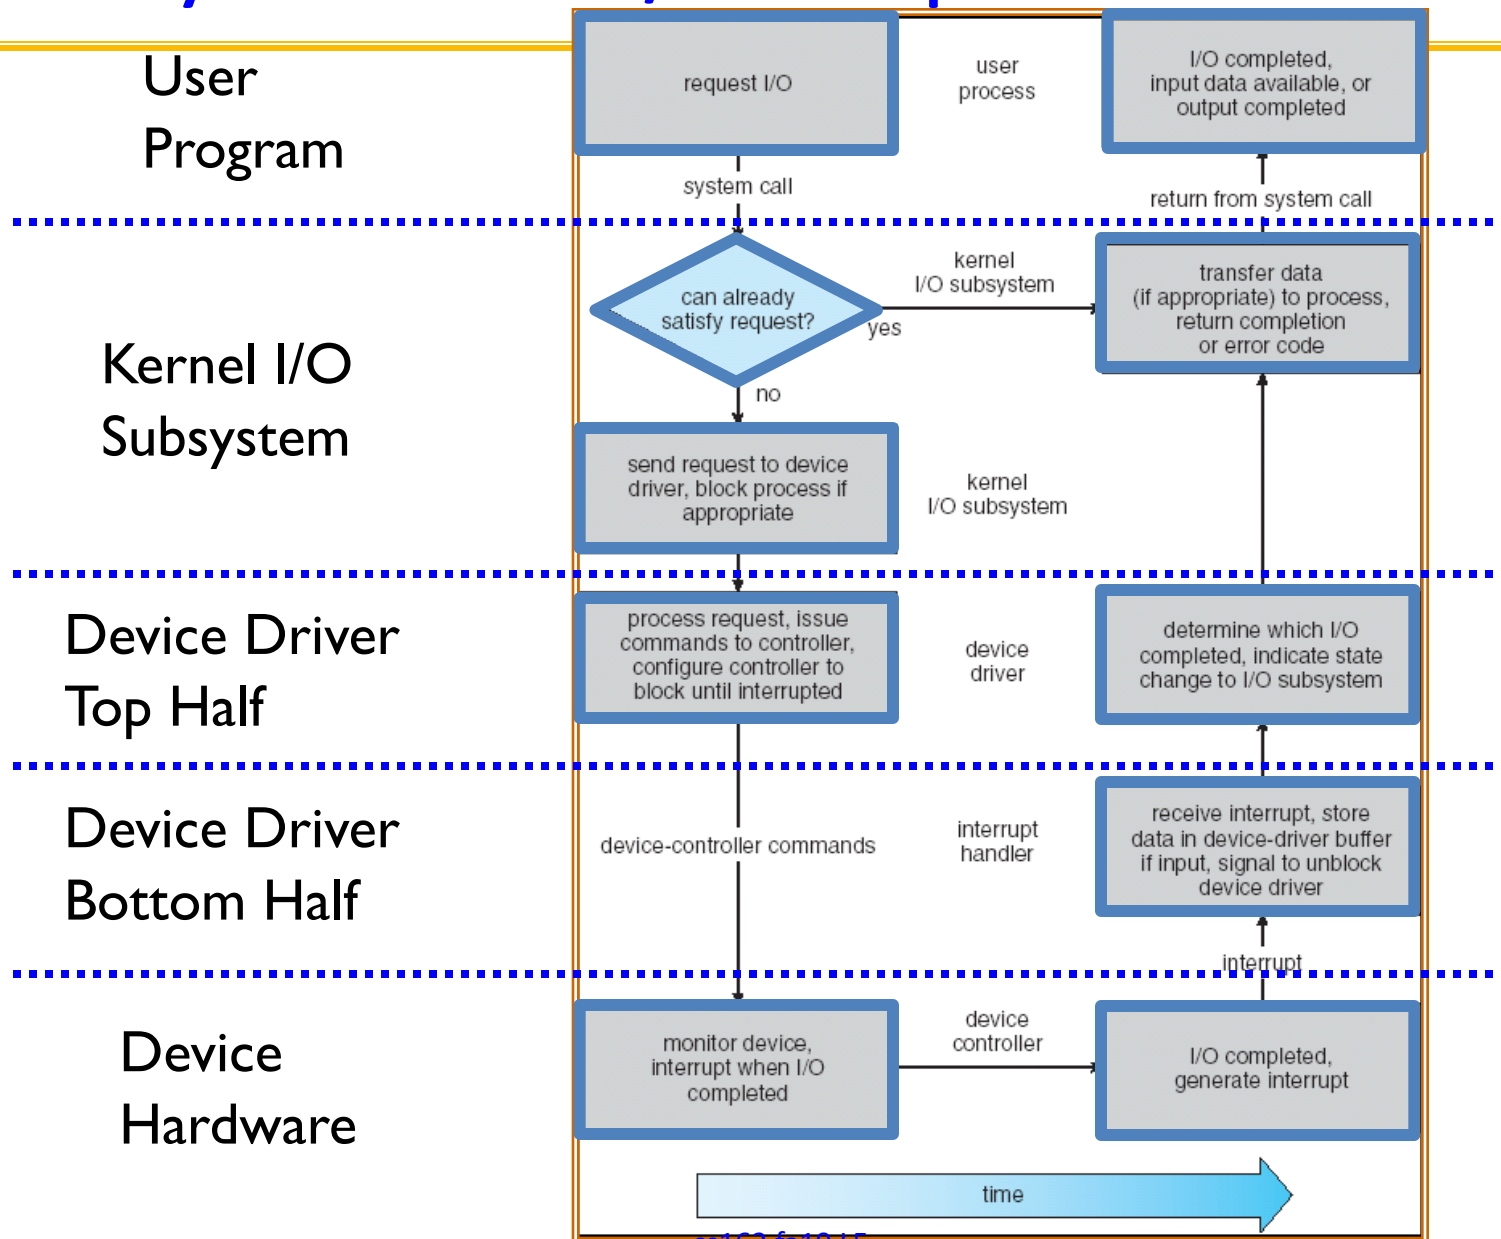
---

- Device-specific code in the kernel that interacts directly with the device hardware
  - Supports a standard, internal interface
  - Same kernel I/O system can interact easily with different device drivers
  - Special device-specific configuration supported with the `ioctl()` system call
- Device Drivers typically divided into two pieces:
  - Top half: accessed in call path from system calls
    - implements a set of **standard, cross-device calls** like `open()`, `close()`, `read()`, `write()`, `ioctl()`, `strategy()`
    - This is the kernel's interface to the device driver
    - Top half will *start* I/O to device, may put thread to sleep until finished
  - Bottom half: run as interrupt routine
    - Gets input or transfers next block of output
    - May wake sleeping threads if I/O now complete





# Life Cycle of An I/O Request





# So what's in our PCB now?

---

- Process ID, name, etc
- Thread object(s) – TCBs
  - Place to save registers when not running
  - Thread status, Links to form lists for scheduling
- Thread Stack
- Lock object for *any lock used by its kernel thread*
  - User level lock info (if multithreaded processes)
- Current working directory
- File Descriptors/Handles for open files



# BIG OS Concepts so far

---

- Processes
- Address Space
- Protection
- Dual Mode
- Interrupt handlers (including syscall and trap)
- Threads
- Synchronization Operations
- File System
  - Integrates processes, users, cwd, protection
- Key Layers: OS Lib, Syscall, Subsystem, Driver
  - User handler on OS descriptors
- Process control
  - fork, wait, signal, exec